

✓ Deep Learning Project

By Victoria Lassner DSML 4220

Goal: Fine tune a model for abstractive Summarization.

Potential models: LLaMa 2 and T5 or GPT

Websites: <https://huggingface.co/docs/transformers/tasks/summarization>

<https://wandb.ai/mostafaibrahim17/ml-articles/reports/Fine-Tuning-LLaMa-2-for-Text-Summarization--Vmlldzo2NjA1OTAy>

<https://wandb.ai/mostafaibrahim17/ml-articles/reports/Crafting-Superior-Summaries-The-ChatGPT-Fine-Tuning-Guide--Vmlldzo1Njc5NDI1>

Definitions:

Abstractive summarization = concise summary of a text by understanding its meaning and creating new sentences, rather than simply extracting phrases from the original text

Extractive Summarization = extracts existing key sentences from passages and creates a new summary.

Potential Datasets: Wikihow: <https://paperswithcode.com/dataset/wikihow> CNN/DailyMail: <https://paperswithcode.com/dataset/cnn-daily-mail-1> BillSum


```
# disables weights and biases
import os
os.environ["WANDB_DISABLED"] = "true"

# downloads packages for model, dataset and tokenizer
# --Quiet limits output of messages
!pip install transformers datasets sentencepiece --quiet

from datasets import load_dataset
from transformers import T5ForConditionalGeneration, TrainingArguments, Trainer, T5Tokenizer
import torch
from torch.utils.data import DataLoader
import torch

# Load CNN/Daily Mail Dataset from dataset package
# will most likely need to downscale again
dataset = load_dataset("cnn_dailymail", "3.0.0")
train_data = dataset["train"].select(range(5000))
```

```
train_data = dataset["train"].select(range(5000))
val_data = dataset["validation"].select(range(500))
```

 /usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning: The secret `HF_TOKEN` does not exist in your Colab secrets. To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>). You will be able to reuse this secret in all of your notebooks. Please note that authentication is recommended but still optional to access public models.
warnings.warn(

```
# preprocess data for model
tokenizer = T5Tokenizer.from_pretrained("t5-small")
```

```
# limit length of input articles and output summary
max_input_length = 512
max_target_length = 150
```

```
def preprocess(examples):
    inputs = ["summarize: " + doc for doc in examples["article"]]
    targets = examples["highlights"]

    model_inputs = tokenizer(
        inputs,
        max_length=max_input_length,
        truncation=True,
        padding="max_length"
    )

    with tokenizer.as_target_tokenizer():
        labels = tokenizer(
            targets,
            max_length=max_target_length,
            truncation=True,
            padding="max_length"
        )

    model_inputs["labels"] = labels["input_ids"]
    return model_inputs
```

```
train_dataset = train_data.map(preprocess, batched=True, remove_columns=["article", "highlights"])
val_dataset = val_data.map(preprocess, batched=True, remove_columns=["article", "highlights"])
```

 tokenizer_config.json: 100% 2.32k/2.32k [00:00<00:00, 32.4kB/s]

spiece.model: 100% 792k/792k [00:00<00:00, 1.49MB/s]

tokenizer.json: 100% 1.39M/1.39M [00:00<00:00, 1.96MB/s]

s]

You are using the default legacy behaviour of the <class 'transformers.models.t5.toke

Map: 100%

5000/5000 [00:29<00:00, 180.89 examples/

<1

```
# Load model T5-small
# T5-base too large without GPU
model = T5ForConditionalGeneration.from_pretrained("t5-small")
```

```
# training arguments
# Will need to scale back to limit training time
# currently just over 8 hours to train
# only for huggingface trainer package
training_args = TrainingArguments(
    output_dir="./t5-cnn-checkpoints",
    per_device_train_batch_size=2,
    per_device_eval_batch_size=2,
    gradient_accumulation_steps=4,
    eval_steps=500,
    save_steps=1000,
    num_train_epochs=2,
    logging_dir='./logs',
    logging_steps=200,
    save_total_limit=2,
    fp16=False,
)
```

```
# for manual training
# batch_size = 2
# learning_rate = 5e-5
# epochs = 2
# gradient_accumulation_steps = 4
# eval_steps = 500
# save_steps = 1000
# logging_steps = 200
# save_total_limit = 2
# fp16 = False
```

config.json: 100%

1.21k/1.21k [00:00<00:00, 85.3kB/s]

Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Fall
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the '

model.safetensors: 100%

242M/242M [00:01<00:00, 179MB/

s]

generation_config.json: 100%

147/147 [00:00<00:00, 5.84kB/

s]

```
# adds padding so shorter sequences match the longest one
```

```

from transformers import DataCollatorForSeq2Seq

data_collator = DataCollatorForSeq2Seq(tokenizer=tokenizer, model=model)

# train model using hugging face's trainer class
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset,
    eval_dataset=val_dataset,
    data_collator=data_collator,
)

trainer.train()

# Porbably will switch to traditional training method

# train_dataloader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
# optimizer = torch.optim.AdamW(model.parameters(), lr=learning_rate)

# model.train()
# for epoch in range(epochs):
#     for step, batch in enumerate(train_dataloader):
#         optimizer.zero_grad()
#         outputs = model(**batch)
#         loss = outputs.loss
#         loss.backward()

#         # Gradient accumulation
#         if (step + 1) % gradient_accumulation_steps == 0:
#             optimizer.step()

#         # Logging
#         if (step + 1) % logging_steps == 0:
#             print(f"Step {step + 1}: Loss = {loss.item()}")

#         # Evaluation
#         if (step + 1) % eval_steps == 0:
#             model.eval()
#             # Run validation logic here
#             model.train()

#     # Save checkpoint manually
#     if (epoch + 1) % save_steps == 0:
#         torch.save(model.state_dict(), f"./t5-cnn-checkpoints/epoch-{epoch+1}.pt")

```

```

-----
NameError                                Traceback (most recent call last)
<ipython-input-1-7a5f6a095494> in <cell line: 0>()
      1 # train model

```

```

----> 2 trainer = Trainer(
      3     model=model,
      4     args=training_args,
      5     train_dataset=train_dataset,

```

NameError: name 'Trainer' is not defined

Next steps: [Explain error](#)

```

#saves current state of model and tokenizer
model.save_pretrained("/content/t5_cnn_model")
tokenizer.save_pretrained("/content/t5_cnn_model")

```

```

-----
NameError                                Traceback (most recent call last)
<ipython-input-2-234c03f2de0c> in <cell line: 0>()
      1 #save model
----> 2 model.save_pretrained("/content/t5_cnn_model")
      3 tokenizer.save_pretrained("/content/t5_cnn_model")

```

NameError: name 'model' is not defined

Next steps: [Explain error](#)

```

# download model from colab for web app to use
from google.colab import files
!zip -r t5_cnn_model.zip /content/t5_cnn_model
files.download("t5_cnn_model.zip")

```

```

adding: content/t5_cnn_model/ (stored 0%)
adding: content/t5_cnn_model/generation_config.json (deflated 29%)
adding: content/t5_cnn_model/special_tokens_map.json (deflated 85%)
adding: content/t5_cnn_model/model.safetensors (deflated 9%)
adding: content/t5_cnn_model/spiece.model (deflated 48%)
adding: content/t5_cnn_model/added_tokens.json (deflated 83%)
adding: content/t5_cnn_model/tokenizer_config.json (deflated 94%)
adding: content/t5_cnn_model/config.json (deflated 63%)

```

```

# test to see how well model is working
def summarize(text):
    input_text = "summarize: " + text
    inputs = tokenizer.encode(input_text, return_tensors="pt", max_length=512, truncation
    summary_ids = model.generate(inputs, max_length=150, min_length=30, length_penalty=2.
    return tokenizer.decode(summary_ids[0], skip_special_tokens=True)

test_article = val_data[0]["article"]

```

```
print("📄 Original Article:\n", test_article[:500], "... \n")  
print("🧠 Summary:\n", summarize(test_article))
```