

## ✓ Deep Learning Project

By Victoria Lassner DSML 4220

**Goal:** Fine tune a model for abstractive Summarization.

**Model:** T5-Base with its Tokenizer

Websites: <https://huggingface.co/docs/transformers/tasks/summarization>

**Future Models to Compare:**

<https://wandb.ai/mostafaibrahim17/ml-articles/reports/Fine-Tuning-LLaMa-2-for-Text-Summarization--Vmlldzo2NjA1OTAy>

<https://wandb.ai/mostafaibrahim17/ml-articles/reports/Crafting-Superior-Summaries-The-ChatGPT-Fine-Tuning-Guide--Vmlldzo1Njc5NDI1>

**Definitions:**

Abstractive summarization = concise summary of a text by understanding its meaning and creating new sentences, rather than simply extracting phrases from the original text.

---

**Dataset:** CNN/DailyMail: <https://paperswithcode.com/dataset/cnn-daily-mail-1> BillSum

```
# disables weights and biases
import os
os.environ["WANDB_DISABLED"] = "true"


# downloads packages for model, dataset and tokenizer
# --Quiet limits output of messages
!pip install transformers datasets evaluate sentencepiece rouge_score --quiet

# Download packages
from datasets import load_dataset, concatenate_datasets
from transformers import T5ForConditionalGeneration, TrainingArguments, Trainer, T5Tokenizer
import torch
from torch.utils.data import DataLoader
import torch

# Load CNN/Daily Mail Dataset from dataset package
# Limit samples to 7000 total.

train_sample_limit = 5000
val_sample_limit = 2000
```

```
dataset = load_dataset("cnn_dailymail", "3.0.0")
limited_train_data = dataset["train"].select(range(train_sample_limit))
limited_val_data = dataset["validation"].select(range(val_sample_limit))
```

 `/usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning`  
The secret `HF\_TOKEN` does not exist in your Colab secrets.  
To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>)  
You will be able to reuse this secret in all of your notebooks.  
Please note that authentication is recommended but still optional to access public models.  
`warnings.warn(`

```
# preprocess data for model
tokenizer = T5Tokenizer.from_pretrained("t5-base")

# limit length of input articles and output summary
max_input_length = 512
max_target_length = 250

chunk_size = 1000

def preprocess(examples):
    inputs = ["summarize: " + doc for doc in examples["article"]]
    targets = examples["highlights"]

    model_inputs = tokenizer(
        inputs,
        max_length=max_input_length,
        truncation=True,
        padding="max_length"
    )

    with tokenizer.as_target_tokenizer():
        labels = tokenizer(
            targets,
            max_length=max_target_length,
            truncation=True,
            padding="max_length"
        )

    # Replace pad token with -100 to ignore in loss
    labels["input_ids"] = [
        [(label if label != tokenizer.pad_token_id else -100) for label in label_seq]
        for label_seq in labels["input_ids"]
    ]

    model_inputs["labels"] = labels["input_ids"]
    return model_inputs
```

```

def process_in_chunks(dataset, chunk_size, preprocess_fn):
    total_len = len(dataset)
    processed_chunks = []

    for i in range(0, total_len, chunk_size):
        chunk = dataset.select(range(i, min(i + chunk_size, total_len)))
        processed_chunk = chunk.map(
            preprocess_fn,
            batched=True,
            remove_columns=["article", "highlights", "id"]
        )
        processed_chunks.append(processed_chunk)

    return concatenate_datasets(processed_chunks)

# Process the training and validation data into chunks
train_dataset = process_in_chunks(limited_train_data, chunk_size, preprocess)
val_dataset = process_in_chunks(limited_val_data, chunk_size, preprocess)

```

You are using the default legacy behaviour of the <class 'transformers.models.t5.tokenizers.tokenizer\_base.T5TokenizerBase'>

Map: 100% 1000/1000 [00:11<00:00, 86.34 examples/s]

/usr/local/lib/python3.11/dist-packages/transformers/tokenization\_utils\_base.py:3980: warnings.warn(

Map: 100% 1000/1000 [00:04<00:00, 226.98 examples/s]

Map: 100% 1000/1000 [00:05<00:00, 188.34 examples/s]

Map: 100% 1000/1000 [00:04<00:00, 223.64 examples/s]

Map: 100% 1000/1000 [00:04<00:00, 225.00 examples/s]

```

import evaluate
import numpy as np

```

```
rouge = evaluate.load("rouge")
```

```

def compute_metrics(eval_preds):
    preds, labels = eval_preds

```

```

    # If preds are logits, convert to token IDs
    if isinstance(preds, tuple):
        preds = preds[0]

```

```

    if preds.ndim == 3: # logits
        preds = np.argmax(preds, axis=-1)

```

```

# Optional safety: clip token IDs to vocab size
preds = np.clip(preds, 0, tokenizer.vocab_size - 1)

decoded_preds = tokenizer.batch_decode(preds, skip_special_tokens=True)
labels = np.where(labels != -100, labels, tokenizer.pad_token_id)
decoded_labels = tokenizer.batch_decode(labels, skip_special_tokens=True)

result = rouge.compute(predictions=decoded_preds, references=decoded_labels, use_stem=True)
return {k: round(v * 100, 2) for k, v in result.items()}

# Load model T5-base
model = T5ForConditionalGeneration.from_pretrained("t5-base")

from transformers import Seq2SeqTrainingArguments

training_args = Seq2SeqTrainingArguments(
    output_dir="./t5-cnn-model",
    eval_steps=500,
    per_device_train_batch_size=4,
    per_device_eval_batch_size=4,
    predict_with_generate=True,
    generation_max_length=128,
    logging_steps=100,
    save_steps=1000,
    num_train_epochs=3,
    fp16=True
)

Using the `WANDB_DISABLED` environment variable is deprecated and will be removed in

# adds padding so shorter sequences match the longest one
from transformers import DataCollatorForSeq2Seq

data_collator = DataCollatorForSeq2Seq(tokenizer=tokenizer, model=model)

# train model using hugging face's trainer class
from transformers import Seq2SeqTrainer

trainer = Seq2SeqTrainer(
    model=model,
    args=training_args,
    train_dataset=train_dataset,
    eval_dataset=val_dataset,
    tokenizer=tokenizer,
    compute_metrics=compute_metrics
)
trainer.train()

```

```
trainer.train()
```

```
<ipython-input-9-dd2b587ad4d3>:4: FutureWarning: `tokenizer` is deprecated and will be removed in a future version of PytorchSeq2SeqWrapper
  trainer = Seq2SeqTrainer(
Passing a tuple of `past_key_values` is deprecated and will be removed in a future version of PytorchSeq2SeqWrapper
[3750/3750 31:31, Epoch 3/3]
```

Step	Training Loss
100	1.734000
200	1.629400
300	1.631200
400	1.662200
500	1.590900
600	1.598300
700	1.618200
800	1.612200
900	1.611600
1000	1.608600
1100	1.558500
1200	1.630200
1300	1.555100
1400	1.458100
1500	1.450800
1600	1.497700
1700	1.511900
1800	1.484700
1900	1.436800
2000	1.489600
2100	1.484800
2200	1.485900
2300	1.448900
2400	1.504600
2500	1.443900
2600	1.416300
----	-----

2700	1.375300
2800	1.379100
2900	1.377600
3000	1.389600
3100	1.349000
3200	1.362700
3300	1.409300
3400	1.452300
3500	1.405300
3600	1.386000
3700	1.403400

```
TrainOutput(global_step=3750, training_loss=1.497725351969401,
metrics={'train_runtime': 1893.5016, 'train_samples_per_second': 7.922,
'train_steps_per_second': 1.98, 'total_flos': 9134368358400000.0, 'train_loss':
1.497725351969401, 'epoch': 3.0})
```

```
metrics = trainer.evaluate()
print(metrics)
```

```
[500/500 15:57]
{'eval_loss': 1.878161907196045, 'eval_rouge1': 37.46, 'eval_rouge2': 16.15, 'eval_rc
```

```
#saves current state of model and tokenzier
model.save_pretrained("/content/t5_cnn_model_base_v4")
tokenizer.save_pretrained("/content/t5_cnn_model_base_v4")

('/content/t5_cnn_model_base_v4/tokenizer_config.json',
'/content/t5_cnn_model_base_v4/special_tokens_map.json',
'/content/t5_cnn_model_base_v4/spiece.model',
'/content/t5_cnn_model_base_v4/added_tokens.json')
```

```
from huggingface_hub import HfApi, HfFolder
from transformers import AutoModelForSequenceClassification, AutoTokenizer
from huggingface_hub import login
```

```
login()
```

```
# Load model and tokenizer
model = T5ForConditionalGeneration.from_pretrained("/content/t5_cnn_model_base_v4")
tokenizer = T5Tokenizer.from_pretrained("/content/t5_cnn_model_base_v4")
```

```
# Save to HuggingFace
```

```
model.push_to_hub("vlassner01/t5_cnn_model_base_v4")
tokenizer.push_to_hub("vlassner01/t5_cnn_model_base_v4")
```

```
olab_upload not found in /root/.cache/huggingface/stored
```

```
model.safetensors: 100% 892M/892M [00:25<00:00, 39.7MB/s]
```

```
README.md: 100% 5.17k/5.17k [00:00<00:00, 505kB/s]
```

```
Commit to HF (commit url: https://huggingface.co/vlassner01/t5\_cnn\_model\_base\_v4/)
```

```
# Step 1: Install and update Hugging Face Hub
```

```
!pip install --upgrade huggingface-hub
```

```
# Step 2: Authenticate to Hugging Face (ensure you enter your API token)
```

```
from huggingface_hub import login
```

```
login() # Enter your Hugging Face API token when prompted
```

```
# Step 3: Prepare the tokenizer files (Ensure all tokenizer files are in this folder)
```

```
!mkdir -p /content/hf_tokenizer_upload
```

```
!cp -r /content/t5_cnn_model_base_v3/* /content/hf_tokenizer_upload/
```

```
# Step 4: Upload tokenizer files to Hugging Face using hf_hub
```

```
from huggingface_hub import upload_file
```

```
# Define your Hugging Face repo name
```

```
repo_name = "vlassner01/t5_cnn_model_base_v4"
```

```
# Path to the local folder containing your tokenizer files
```

```
folder_path = '/content/hf_tokenizer_upload'
```

```
# Upload individual files to Hugging Face
```

```
upload_file(
    path_or_fileobj=f"{folder_path}/spiece.model", # Replace with actual file path
    path_in_repo="spiece.model", # Path in the Hugging Face repo
    repo_id=repo_name,
    commit_message="Upload spiece.model"
)
```

```
upload_file(
    path_or_fileobj=f"{folder_path}/tokenizer_config.json", # Replace with actual file p
    path_in_repo="tokenizer_config.json", # Path in the Hugging Face repo
    repo_id=repo_name,
    commit_message="Upload tokenizer_config.json"
)
```

```
upload_file(
    path_or_fileobj=f"{folder_path}/special_tokens_map.json", # Replace with actual file
    path_in_repo="special_tokens_map.json", # Path in the Hugging Face repo
```

```

repo_id=repo_name,
commit_message="Upload special_tokens_map.json"
)

upload_file(
    path_or_fileobj=f"{folder_path}/tokenizer.json", # Replace with actual file path
    path_in_repo="tokenizer.json", # Path in the Hugging Face repo
    repo_id=repo_name,
    commit_message="Upload tokenizer.json"
)

# Step 5: Verify the files
# Once the upload finishes, check your model page at:
# https://huggingface.co/vlassner01/t5_cnn_model_base_v3

```

```

Requirement already satisfied: huggingface-hub in /usr/local/lib/python3.11/dist-pack
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (f
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.11/dist-pac
Requirement already satisfied: packaging>=20.9 in /usr/local/lib/python3.11/dist-pack
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (f
Requirement already satisfied: tqdm>=4.42.1 in /usr/local/lib/python3.11/dist-package
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.1
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-package
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-p
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-p
olab_upload not found in /root/.cache/huggingface/stored

```

```

spiece.model: 100% 792k/792k [00:00<00:00, 4.55MB/

```

```

s]

```

```

No files have been modified since last commit. Skipping to prevent empty commit.
WARNING:huggingface_hub.hf_api:No files have been modified since last commit. Skippir
No files have been modified since last commit. Skipping to prevent empty commit.
WARNING:huggingface_hub.hf_api:No files have been modified since last commit. Skippir

```

```

-----
ValueError                                Traceback (most recent call last)
<ipython-input-16-805077f8d776> in <cell line: 0>()
    41 )
    42
--> 43 upload_file(
    44     path_or_fileobj=f"{folder_path}/tokenizer.json", # Replace with actual
file path
    45     path_in_repo="tokenizer.json", # Path in the Hugging Face repo

```

4 frames

```

/usr/local/lib/python3.11/dist-packages/huggingface_hub/_commit_api.py in
__post_init__(self)
    180         path_or_fileobj =
os.path.normpath(os.path.expanduser(self.path_or_fileobj))
    181         if not os.path.isfile(path_or_fileobj):

```



```
181         if not os.path.isfile(path_or_fileobj):  
--> 182             raise ValueError(f"Provided path: '{path_or_fileobj}' is not  
a file on the local file system")
```